

## 1. L1 and L2 Loss Functions with Benefits and Tradeoffs

### L2 Loss (Ridge Regression)

- **Objective**:  $\min_{\beta} \|X\beta - y\|_2^2 + \lambda \|\beta\|_2^2$  - **Closed-form solution**:  $\beta^* = (X^T X + \lambda I)^{-1} X^T y$  - **Benefits**: Shrinks coefficients, prevents overfitting, retains all features. - **Tradeoff**: Does not perform feature selection, which reduces interpretability in high-dimensional spaces.

### L1 Loss (Lasso Regression)

- **Objective**:  $\min_{\beta} \|X\beta - y\|_2^2 + \lambda \|\beta\|_1$  - **No closed-form solution**: Requires iterative optimization (e.g., coordinate descent). - **Benefits**: Performs feature selection by driving some coefficients to zero, improving interpretability. - **Tradeoff**: Excludes features that may still hold important information, may underperform if many features are relevant.

## 2. Ridge Regularized Least Squares

- **Objective**:  $\|X\beta - y\|_2^2 + \lambda \|\beta\|_2^2$  - **Closed-form solution**:  $\beta^* = (X^T X + \lambda I)^{-1} X^T y$  - **Benefits**: Regularization avoids overfitting by shrinking coefficients. - **Tradeoff**: Requires careful tuning of  $\lambda$  for the right balance between bias and variance.

## 3. Loss Functions

### Mean Squared Error (MSE)

- **Objective**:  $\min_{\beta} \frac{1}{n} \sum_{i=1}^n (y_i - X_i \beta)^2$  - **Closed-form solution**:  $\beta^* = (X^T X)^{-1} X^T y$  - **Benefits**: Simple and easy to compute. Works well with linear regression. - **Tradeoff**: Sensitive to outliers, which can dominate the loss.

### Cross-Entropy Loss (Logistic Regression)

- **Objective**:  $L(\beta) = -\sum_{i=1}^n (y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i))$  - **No closed-form solution**: Solved iteratively via gradient descent. - **Benefits**: Ideal for binary classification, models probabilities effectively. - **Tradeoff**: More computationally expensive, requires careful parameter tuning.

### Hinge Loss (Support Vector Machines)

- **Objective**:  $L = \max(0, 1 - y_i X_i \beta)$  - **No closed-form solution**: Solved via convex optimization methods (e.g., quadratic programming). - **Benefits**: Useful in maximizing the margin between classes. - **Tradeoff**: Computationally intensive for large datasets.

### 0-1 Loss (Classification Problems)

- **Objective**:  $L = \sum_{i=1}^n \mathbb{I}(y_i \neq \hat{y}_i)$ , where  $\mathbb{I}$  is an indicator function. - **Benefits**: Easy to understand and directly penalizes misclassification. - **Tradeoff**: Non-convex and discontinuous, so not suitable for gradient-based methods.

### Multi-Class Cross-Entropy Loss

- **Objective**: Used for multi-class classification problems to measure the performance of a classification model whose output is a probability distribution across multiple classes. - **Formula**:  $L = -\sum_{i=1}^n \sum_{c=1}^C y_{i,c} \log(\hat{y}_{i,c})$ , where  $C$  is the number of classes,  $y_{i,c}$  is the true label (1 if class  $c$  is correct, 0 otherwise), and  $\hat{y}_{i,c}$  is the predicted probability for class  $c$ . - **Benefits**: Ideal for problems where each instance can belong to one of several classes (e.g., image classification). - **Models** probabilistic outcomes effectively, providing confidence scores. - **Tradeoffs**: - More computationally intensive compared to binary cross-entropy due to multiple classes. - Sensitive to class imbalance, which may lead to biased predictions if one class dominates. - **Key Concepts**: This loss encourages models to output probabilities that are as close as possible to the true one-hot encoded labels.

### Binary Cross-Entropy Loss

- **Objective**: Used for binary classification problems, measuring the performance of a classification model whose output is a probability value between 0 and 1. - **Formula**:  $L = -\frac{1}{n} \sum_{i=1}^n (y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i))$ , where  $y_i$  is the true binary label (1 or 0), and  $\hat{y}_i$  is the predicted probability for label 1. - **Benefits**: Ideal for binary classification tasks like spam detection or medical diagnosis. - **Tradeoffs**: - Can struggle with class imbalance - Sensitive to extreme predictions (very close to 0 or 1) that may cause large gradients, impacting training stability. - **Key Concepts**: This loss penalizes incorrect predictions and emphasizes confidence, making it widely used in classification problems involving two outcomes.

### Equivalence of Multi-Class and Binary Cross-Entropy Loss

- **Equivalence**: Multi-class cross-entropy simplifies to binary cross-entropy when the number of classes  $C = 2$ . - **Setup**: For binary classification, we set  $\beta^{(0)} = -\beta$  and  $\beta^{(1)} = \beta$ . - **Multi-Class Cross-Entropy** for two classes:

$$L = -\sum_{i=1}^n \sum_{c=0}^1 y_{i,c} \log(\hat{y}_{i,c})$$

where  $\hat{y}_{i,0} = \sigma(-\beta^T x_i)$  and  $\hat{y}_{i,1} = \sigma(\beta^T x_i)$ . - **Simplification**: Plugging in the values of  $\hat{y}_{i,0}$  and  $\hat{y}_{i,1}$ :

$$L = -\sum_{i=1}^n \left( y_i \log(\sigma(\beta^T x_i)) + (1 - y_i) \log(1 - \sigma(\beta^T x_i)) \right)$$

This is the **Binary Cross-Entropy Loss**:

$$L = -\frac{1}{n} \sum_{i=1}^n (y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i))$$

- **Conclusion**: Multi-class cross-entropy for two classes reduces to binary cross-entropy when  $\beta^{(0)} = -\beta$  and  $\beta^{(1)} = \beta$ .

## 4. Gaussian Naive Bayes

- **MAP**:  $\arg\max_y \left( p(y|x) = \frac{p(x|y)p(y)}{p(x)} \right)$  - **Likelihood**:  $p(x|y) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x-\mu)^2}{2\sigma^2}\right)$  - **Benefits**: Fast to compute, assumes independence between features. - **Tradeoff**: Assumption of independence is often unrealistic, which can lead to inaccuracies.

## 5. K-Fold Cross Validation

- **Benefits**: Provides better estimates of model performance by using every data point for both training and validation. - **Tradeoff**: Computationally expensive, especially for large datasets or complex models.

## 6. Derivatives for Optimization

- **Gradient**:  $\nabla f(x) = \left[ \frac{\partial f}{\partial x_1}, \frac{\partial f}{\partial x_2}, \dots, \frac{\partial f}{\partial x_n} \right]$  - **Chain Rule**:  $\frac{\partial f}{\partial x} = \frac{\partial f}{\partial u} \cdot \frac{\partial u}{\partial x}$

## 7. Maximum Likelihood Estimation (MLE)

- **Gaussian MLE**:  $\mu_{MLE} = \frac{1}{n} \sum x_i$ ,  $\sigma_{MLE}^2 = \frac{1}{n} \sum (x_i - \mu)^2$  - **Bernoulli MLE**:  $\mu_{MLE} = \frac{1}{n} \sum x_i$  - **Benefits**: Provides efficient estimators if the assumptions about data distribution are correct. - **Tradeoff**: Assumptions about data distribution can lead to poor results if incorrect.

## 8. Gradient Descent

- **Update Rule**:  $\beta \leftarrow \beta - \alpha \nabla L(\beta)$ , where  $\alpha$  is the learning rate. - **Benefits**: Works for large models without closed-form solutions (e.g., neural networks, logistic regression). - **Tradeoff**: Sensitive to choice of learning rate, can converge slowly or diverge.

## 9. Rayleigh Distribution MLE

- **PDF**:  $p(x) = \frac{x}{\sigma^2} \exp\left(-\frac{x^2}{2\sigma^2}\right)$  - **MLE for  $\sigma$** :  $\sigma_{MLE} = \sqrt{\frac{1}{2n} \sum x_i^2}$  - **Benefits**: Provides a simple estimation method for certain non-negative data. - **Tradeoff**: Assumes a specific distribution, may not generalize well to other data.

## 10. Matrix Calculus Rules

- **Quadratic Form Derivative**:  $\frac{d}{d\beta} (\beta^T X \beta) = 2X\beta$  - **Logarithmic Derivative**:  $\frac{d}{d\beta} \log f(\beta) = \frac{1}{f(\beta)} \cdot f'(\beta)$

## 11. Bias-Variance Tradeoff

- **Benefits**: Helps in understanding model complexity, assisting in selecting simpler models to reduce variance or more complex models to reduce bias. - **Tradeoff**: High bias leads to underfitting (poor accuracy), high variance leads to overfitting (poor generalization).

## 12. Regularization Techniques

### L2 Regularization (Ridge)

- **Objective**: Adds  $\lambda \|\beta\|_2^2$  to the loss function. - **Closed-form solution**:  $\beta^* = (X^T X + \lambda I)^{-1} X^T y$  - **Benefits**: Prevents overfitting, improves generalizability. - **Tradeoff**: Does not eliminate features, making models harder to interpret in high dimensions.

### L1 Regularization (Lasso)

- **Objective**: Adds  $\lambda \|\beta\|_1$  to the loss function. - **No closed-form solution**: Solved via optimization (e.g., coordinate descent). - **Benefits**: Encourages sparsity, making the model more interpretable. - **Tradeoff**: Can exclude relevant features if not tuned carefully.

## 13. One-Hot Encoding

- **Definition**: One-hot encoding is a process used to convert categorical data into a binary vector for each category. - **Process**: Each category in the dataset is transformed into a vector where only one element is 1, and the rest are 0s. - **Benefits**: Allows categorical data to be used in machine learning algorithms that require numerical input. - **Tradeoff**: Can lead to high-dimensional datasets when the number of categories is large, which may increase computational costs and memory usage.

## 14. Distributions

- **Laplace PDF**:  $p(x) = \frac{1}{2b} \exp\left(-\frac{|x-\mu|}{b}\right)$  - **Gaussian PDF**:  $p(x) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x-\mu)^2}{2\sigma^2}\right)$  - **Bernoulli PMF**:  $p(x) = \mu^x (1 - \mu)^{1-x}$  - **Rayleigh PDF**:  $p(x) = \frac{x}{\sigma^2} \exp\left(-\frac{x^2}{2\sigma^2}\right)$